


# Supervised and unsupervised learning

Supervised learning: input data  $\rightarrow$  output data.

Unsupervised learning: input data and nothing else.

# Clustering problem



This problem is very hard to formalize. What exactly is a *cluster*? How to determine the number of clusters in the data? Are the data clustered at all? How to compare two clustering algorithms?

# $K$ -means clustering

$K$ -means clustering aims to cluster the dataset  $\{\mathbf{x}_i\}$  into  $K$  clusters  $\{S_k\}$ , each represented by a vector  $\boldsymbol{\mu}_k$ , to minimize the following loss function:


$$\mathcal{L} = \sum_{k=1}^K \sum_{i \in S_k} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2.$$

Alternatively, this can be written as

$$\mathcal{L} = \sum_{k=1}^K \sum_{i=1}^n r_{ik} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2,$$

where  $r_{ik} = 1$  if  $\mathbf{x}_i \in S_k$  and 0 otherwise.

# $K$ -means loss function



# Minimizing the $K$ -means loss

$$\mathcal{L} = \sum_k \sum_{i \in S_k} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2 = \sum_{k=1}^K \sum_{i=1}^n r_{ik} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2$$

Not analytically solvable. Not convex. Gradient descent can be messy.

Alternative approach (Lloyd's algorithm): iteratively optimize over  $r_{ik}$  and over  $\boldsymbol{\mu}_k$ .


- For fixed  $\boldsymbol{\mu}_k$ : assign each point  $\mathbf{x}_i$  to the nearest cluster center.

$$i \in S_k \text{ if } k = \arg \min_j \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2.$$

- For fixed  $r_{ik}$ :


$$\boldsymbol{\mu}_k = \frac{1}{|S_k|} \sum_{i \in S_k} \mathbf{x}_i.$$

# Illustration of the Lloyd's algorithm







Bishop, *Pattern Recognition and Machine Learning*

# Local minima



<https://stats.stackexchange.com/questions/133656>

# Drawbacks of $K$ -means





# Gaussian mixture model (GMM)

Gaussian mixture:

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k).$$

Intuitively, we could use the same iterative approach as in the Lloyd's algorithm for  $K$ -means:

- Assign each point to the 'nearest' Gaussian component (cluster).  
Here 'nearest' means 'with the highest posterior'  $\pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ .
- Update the parameters  $(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, \pi_k)$  of each Gaussian.

# Likelihood in GMM

Gaussian mixture:

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k).$$

Log-likelihood:

$$\mathcal{L} = \sum_{i=1}^n \log \left[ \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right],$$

where  $\mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) = \dots \exp \left( -\frac{1}{2}(\mathbf{x}_i - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1}(\mathbf{x}_i - \boldsymbol{\mu}_k) \right)$ .

Set the derivative with respect to  $\boldsymbol{\mu}_k$  to zero:

$$\sum_{i=1}^n \underbrace{\frac{\pi_k \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}}_{z_{ik}} \cancel{\boldsymbol{\Sigma}_k^{-1}} (\mathbf{x}_i - \boldsymbol{\mu}_k) = 0.$$

# Likelihood in GMM

Set the derivative with respect to  $\boldsymbol{\mu}_k$  to zero:

$$\sum_{i=1}^n \frac{\pi_k \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\underbrace{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}_{z_{ik}}} \cancel{\boldsymbol{\Sigma}_k^{-1}} (\mathbf{x}_i - \boldsymbol{\mu}_k) = 0.$$

$$\sum_{i=1}^n z_{ik} (\mathbf{x}_i - \boldsymbol{\mu}_k) = 0.$$

$$\boldsymbol{\mu}_k = \frac{\sum z_{ik} \mathbf{x}_i}{\sum z_{ik}}.$$

This is a *weighted* mean of all points.

Very similar derivation shows that  $\boldsymbol{\Sigma}_k$  should be the weighted covariance matrix, and  $\pi_k = \sum z_{ik}/n$ .


# Expectation-maximization (EM)

Expectation-maximization algorithm iteratively alternates between updating  $\mu_k, \Sigma_k, \pi_k$  and updating  $z_{ik}$ :

- **E-step:** compute the posterior probability  $z_{ik}$  for each point to be in each Gaussian component.
- **M-step:** update the parameters ( $\mu_k, \Sigma_k, \pi_k$ ) of each Gaussian using weighted averages.

EM is a very generic algorithm to optimize likelihood in probabilistic models with *latent variables*. (In GMMs, latent variables are true class memberships.) E-step computes posterior over latent variables, conditioned on the parameters of the model. M-step optimizes the parameters, conditioned on the latent variables.

# Illustration of the EM




Bishop, *Pattern Recognition and Machine Learning*

# Divergence in GMM

Gaussian mixture:

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k).$$

The likelihood can diverge if  $\boldsymbol{\mu}_k = \mathbf{x}_i$  for some  $i$  and  $\boldsymbol{\Sigma}_k \rightarrow \mathbf{0}$ .



In practice: if one of the Gaussians starts ‘collapsing’ during EM towards a degenerate solution, do something (e.g. randomly reset its mean and covariance matrix).

# EM vs. gradient descent

- Both EM and gradient descent are iterative algorithms.
- Both can converge to a local minimum.
- EM does not need a learning rate.
- In EM, all parameters are automatically meaningful after each step without imposing constraints (such as  $\pi_k$  summing to 1, or all  $\Sigma_k$  being positive-definite).

# GMM vs. $K$ -means

Similar to what we discussed about LDA, one can constrain  $\Sigma_k$  in a GMM to be shared between classes, or diagonal, or spherical.

A GMM with shared spherical covariance matrix  $\Sigma_k = \sigma^2 \mathbf{I}$  is very closely related to  $K$ -means. The main difference is that  $K$ -means performs *hard* cluster assignments in the 'E-step', whereas GMM performs *soft* cluster assignments. If  $\sigma^2 \rightarrow 0$ , GMM converges to  $K$ -means.

Note: in practical implementations it can be convenient to initialize GMM with a  $K$ -means solution.



# GMM vs. $K$ -means

